Sentinel - Demo and Evaluation Guide

How to run the platform and what to show in a 10-minute walkthrough

Prepared for incubation evaluation | 2026 | Live deployment: http://localhost:8000

1. Prerequisites

- Python 3.11+ with backend requirements installed (backend/requirements.txt).
- Node 18+; frontend at frontend/ builds with tsc + Vite into frontend/dist (served by the backend).
- Ollama running locally with the `sentinel` (or `mistral`) model pulled - provides offline AI inference.

2. Run the Stack

- Backend: python -m uvicorn app.main:app --host 127.0.0.1 --port 8000 (from backend/).
- Frontend: already built to frontend/dist and served at http://localhost:8000 by the backend.
- Rebuild frontend after changes: node node_modules/typescript/bin/tsc -b then node node_modules/vite/bin/vite.js build (from frontend/).
- Seed demo data: python seed_demo.py (from backend/).

3. Demo Accounts

Role	Username	Password
Psychologist	cel	1234
Patient	alaya	4321

4. 10-Minute Walkthrough Script

4.1 Patient portal (2 min)

- Log in as alaya. Show biometric trends (BPM, stress, sleep, SpO2, mood) on the dashboard.
- Open an AI journal summary; note the 28-label emotion classification feeding the summary.

4.2 Psychologist triage (3 min)

- Log in as cel. Open the Priority Triage Dashboard: patients are ranked by risk score; expand a patient to see the five bio metric cards, the AI clinical insight, and the explainability panel.
- Trigger the 'Why this summary?' panel to show the traceable, non-black-box reasoning.

4.3 Crisis escalation (3 min)

- From the patient view, trigger a crisis (or POST /api/crisis/trigger). Watch it appear live on the psychologist dashboard, then acknowledge it to freeze the escalation timer.
- Explain the staged protocol: patient siren, trusted-contact signed email, helpline escalation, halt-on-ack.

4.4 Validation artifacts (2 min)

- Show backend/benchmarks/logbook_benchmark.csv (47 timed runs) and quote the headline numbers: 96% discrepancy accuracy at 0.1 ms, 98/98 tests, 1.4-1.7 s local AI, 7-41 ms REST.

- Reproduce live: cd backend && python -m benchmarks.runner && python -m pytest tests -q.

5. What Evaluators Can Check Themselves

- Security: bcrypt policy, JWT rotation, Fernet field encryption, hash-chained audit log, signed links.
- Privacy: the local-AI path sends no data off-device; cloud AI is off by default.
- Openness: full source in the repository; benchmark and test harness committed; docs regenerate from scripts in the root.

Companion documents: White Paper (03), Validation & Timing Report (04), Technical Design (05), Deployment & Setup Guide (10).